

Complete GCP Deployment Workflow Guide · Generated February 19, 2026

## Table of Contents

---

1. [GCP Services Used & Their Purpose](#)
  2. [One-Time Deployment Workflow \(Step-by-Step\)](#)
  3. [Runtime Workflow \(When Cloud Scheduler Fires\)](#)
  4. [Visual Diagrams](#)
- 

## 1. GCP Services Used & Their Purpose

---

#	GCP Service	Purpose in This Project	Used During
1	Cloud Build	Builds the Docker image from the Dockerfile on Google's servers	Deployment only
2	Artifact Registry / Container Registry	Stores the built Docker image ( <code>gcr.io/{PROJECT_ID}/sap-doc-monitor</code> )	Deployment (image storage for Cloud Run)
3	Cloud Run	Runs the containerized Python app (Flask HTTP server) that performs the actual monitoring	Deployment + Runtime
4	Cloud Storage (GCS)	Persistently stores document snapshots ( <code>.txt</code> files) between runs — because Cloud Run containers are <b>ephemeral</b> (destroyed after each run)	Runtime only
5	Secret Manager	Securely stores the email SMTP password ( <code>email-password</code> ). Cloud Run injects it as the <code>EMAIL_PASSWORD</code> environment variable at runtime	Runtime only
6	Service Account ( <code>sap-monitor-scheduler</code> )	An identity that gives Cloud Scheduler permission to invoke the <b>private</b> (unauthenticated access blocked) Cloud Run service using OIDC tokens	Runtime only
7	Cloud Scheduler	The <b>trigger</b> — sends an HTTP POST request to Cloud Run on a cron schedule (twice daily at 9 AM and 6 PM) to start the monitoring job	Runtime only

## Key Clarification

*GCS, Secret Manager, Service Account, and Cloud Scheduler are NOT part of the Docker build/push/deploy process. They are runtime infrastructure — they exist so the app can function correctly every time Cloud Run executes it.*

## 2. One-Time Deployment Workflow (Step-by-Step)

Below is the **exact sequence** from the deployment script. Each step must complete before the next one starts.

### Step 1: Configure gcloud CLI

```
gcloud config set project {PROJECT_ID}
gcloud config set run/region us-central1
```

**What happens:** Sets your GCP project and region so all subsequent commands target the correct project.

## Step 2: Enable Required GCP APIs

```
gcloud services enable \
run.googleapis.com \
cloudscheduler.googleapis.com \
cloudbuild.googleapis.com \
storage.googleapis.com \
secretmanager.googleapis.com
```

**What happens:** Activates 5 GCP APIs. Without this, none of the services can be used.

API	Enables
run.googleapis.com	Cloud Run
cloudscheduler.googleapis.com	Cloud Scheduler
cloudbuild.googleapis.com	Cloud Build
storage.googleapis.com	Cloud Storage (GCS)
secretmanager.googleapis.com	Secret Manager

## Step 3: Create Cloud Storage (GCS) Bucket

```
gsutil mb -p {PROJECT_ID} -l us-central1 gs://-{PROJECT_ID}-sap-snapshots
```

**What happens:** Creates a GCS bucket named `{PROJECT_ID}-sap-snapshots`.

**Why this is needed:** - Cloud Run containers are **ephemeral** — all files inside them are destroyed when the container shuts down. - The app needs to compare **current** page content with **previous** snapshots to detect changes. - GCS provides **persistent storage** that survives across container restarts. - **Before each run:** All previous snapshot `.txt` files are **downloaded** from GCS into the container's local `snapshots/` directory. These downloaded snapshots represent the "last known state" of each SAP documentation page. The app then fetches the **live** content from SAP Help, extracts text, and compares it line-by-line against the downloaded snapshot for each page. If differences are found (additions, removals, or structural changes), those are flagged as changes. -

**After each run:** The updated snapshots (with the latest content) are **uploaded** back to GCS, so the **next** scheduled run can download them and repeat the comparison cycle.

Also grants the App Engine default service account `objectAdmin` permission on the bucket:

```
gsutil iam ch serviceAccount:{PROJECT_ID}@appspot.gserviceaccount.com:objectAdmin gs://{BUCKET_NAME}
```

**Why this permission is needed:** - Cloud Run runs your container under the **App Engine default service account** (`{PROJECT_ID}@appspot.gserviceaccount.com`). - The app code calls GCS APIs to download and upload snapshot files (`blob.download_to_filename()`, `blob.upload_from_filename()`, `blob.delete()`). - These API calls are authenticated as whatever service account Cloud Run is running under. - Without `objectAdmin` on the bucket, these calls would fail with **403 Permission Denied**. - `objectAdmin` grants full read/write/delete access to objects in the bucket — exactly what the app needs to download old snapshots, upload updated ones, and delete stale ones.

## Step 4: Create Secret in Secret Manager

```
echo -n "{EMAIL_PASSWORD}" | gcloud secrets create email-password \
--data-file=--replication-policy="automatic"
```

**What happens:** Stores the email SMTP password as a secret named `email-password` in Secret Manager.

**Why this is needed:** - The app sends email notifications via SMTP (Office 365). - Hardcoding passwords in code or environment variables is insecure. - Secret Manager encrypts and securely stores the password. - Cloud Run mounts it at runtime as the `EMAIL_PASSWORD` environment variable (via `--update-secrets` flag during deploy).

Also grants the service account permission to access the secret:

```
gcloud secrets add-iam-policy-binding email-password \
--member="serviceAccount:{PROJECT_ID}@appspot.gserviceaccount.com" \
--role="roles/secretmanager.secretAccessor"
```

## Step 5: Build Docker Image (Cloud Build)

```
gcloud builds submit --tag gcr.io/{PROJECT_ID}/sap-doc-monitor
```

**What happens (inside Google Cloud):** 1. Your local source code (Dockerfile + `sap-doc-monitor/` folder) is uploaded to Cloud Build. 2. Cloud Build reads the `Dockerfile` and builds the image: - Starts from `python:3.11-slim` - Installs Chrome browser + Selenium dependencies - Installs Python packages from

`requirements.txt` - Copies application code - Replaces `settings.py` with `settings.cloud.py` (reads from env vars instead of hardcoded values) - Sets the entrypoint to `python cloud_run_app.py` (Flask HTTP server) 3. The built image is automatically pushed to Container Registry at `gcr.io/{PROJECT_ID}/sap-doc-monitor`.

*This is the "build & push" step you already know. Cloud Build does both in a single command — you don't need to run `docker build` and `docker push` separately.*

## Step 6: Deploy to Cloud Run

```
gcloud run deploy sap-doc-monitor \
  --image gcr.io/{PROJECT_ID}/sap-doc-monitor \
  --platform managed \
  --region us-central1 \
  --no-allow-unauthenticated \
  --memory 2Gi \
  --cpu 2 \
  --timeout 900 \
  --max-instances 1 \
  --set-env-vars EMAIL_SENDER="..." \
  --set-env-vars EMAIL_RECEIVER="..." \
  --set-env-vars SMTP_SERVER="..." \
  --set-env-vars SMTP_PORT="587" \
  --set-env-vars BASE_DOCUMENTATION_URL="..." \
  --set-env-vars SNAPSHOT_DIR="/app/snapshots" \
  --update-secrets EMAIL_PASSWORD=email-password:latest
```

**What happens:** 1. Cloud Run pulls the Docker image from Container Registry. 2. Creates a managed Cloud Run service named `sap-doc-monitor`. 3. Configures it with: - `--no-allow-unauthenticated` — Only authenticated requests (with valid OIDC tokens) can invoke the service. Public access is blocked. - `--memory 2Gi --cpu 2` — Sufficient resources for running headless Chrome. - `--timeout 900` — 15-minute max execution time (scraping all pages takes time). - `--max-instances 1` — Only one container runs at a time (prevents duplicate runs). - `--set-env-vars` — Injects email config and doc URL as environment variables. - `--update-secrets EMAIL_PASSWORD=email-password:latest` — Mounts the Secret Manager secret as the `EMAIL_PASSWORD` env var. 4. Returns a Service URL (e.g., `https://sap-doc-monitor-xxxxx-uc.a.run.app`).

*The Flask app (`cloud_run_app.py`) listens on port 8080. When it receives a POST request at `/`, it calls `main.main()` which runs the full monitoring workflow.*

## Step 7: Create Service Account for Cloud Scheduler

```
gcloud iam service-accounts create sap-monitor-scheduler \
--display-name "SAP Monitor Scheduler"
```

Then grant it permission to invoke the Cloud Run service:

```
gcloud run services add-iam-policy-binding sap-doc-monitor \
--member="serviceAccount:sap-monitor-scheduler@{PROJECT_ID}.iam.gserviceaccount.com" \
--role="roles/run.invoker" \
--region=us-central1
```

**Why this is needed:** - In Step 6, the Cloud Run service was deployed with `--no-allow-unauthenticated` (private). - Cloud Scheduler needs an **identity** to authenticate its requests to Cloud Run. - This service account has the `roles/run.invoker` role, which allows it to call the Cloud Run service. - Cloud Scheduler uses this service account to generate **OIDC tokens** attached to each HTTP request.

## Step 8: Create Cloud Scheduler Job

```
gcloud scheduler jobs create http sap-doc-monitor-job \
--location=us-central1 \
--schedule="0 9,18 * * *" \
--uri={CLOUD_RUN_SERVICE_URL} \
--http-method=POST \
--oidc-service-account-email=sap-monitor-scheduler@{PROJECT_ID}.iam.gserviceaccount.com \
--oidc-token-audience={CLOUD_RUN_SERVICE_URL} \
--time-zone="America/New_York"
```

**What happens:** Creates a Cloud Scheduler job that: - Runs on a **cron schedule** (`0 9,18 * * *` = every day at 9:00 AM and 6:00 PM). - Sends an **HTTP POST** to the Cloud Run service URL. - Authenticates using **OIDC token** signed as the `sap-monitor-scheduler` service account.

*This is the trigger that makes the entire system automated. Without Cloud Scheduler, you would have to manually call the Cloud Run URL every time.*

## Step 9: Test the Deployment

```
gcloud scheduler jobs run sap-doc-monitor-job --location=us-central1
```

**What happens:** Manually triggers the scheduler job to verify the full pipeline works end-to-end.

### 3. Runtime Workflow (When Cloud Scheduler Fires)

Every time Cloud Scheduler triggers (daily at 9 AM and 6 PM, or manually), here is the **exact sequence of events**:

#### Stage 1: Cloud Scheduler → Cloud Run (THE TRIGGER)

```
Cloud Scheduler fires at 9:00 AM or 6:00 PM
|
|--- Generates an OIDC token (signed as sap-monitor-scheduler service account)
|--- Sends HTTP POST to Cloud Run service URL
|
|↓
Cloud Run receives the authenticated HTTP POST request
|
|--- Verifies OIDC token → authenticated ✓
|--- Spins up a fresh container (cold start if no warm instance)
|--- Flask app (cloud_run_app.py) handles request at route '/'
|--- Calls main.main() – starts the monitoring logic
|
|↓
```

#### Stage 2: Download Previous Snapshots from GCS

```
main.main() starts
|
|--- Checks: is GCS_BUCKET_NAME env var set? (is_gcs_enabled())
|
|--- YES → GCS Mode (Cloud Run):
|   |--- Wipes any local .txt snapshots baked into the Docker image
|   |--- Downloads ALL previous .txt snapshot files from GCS bucket
|   |   |--- (gs://{{PROJECT_ID}}-sap-snapshots/snapshots/*.txt)
|   |--- These represent the "last known state" of each SAP doc page
|
|--- NO → Local Mode (development):
|   |--- Uses snapshots already in the local snapshots/ directory
|
|↓
```

**Why this stage exists:** Cloud Run containers are ephemeral — every container starts fresh with no memory of previous runs. GCS acts as the "persistent memory" between runs.

### Stage 3: Discover & Fetch SAP Documentation Pages

- Auto-discovers all documentation page URLs from SAP Help TOC  
(Uses Selenium + headless Chrome to load the TOC page)
- For each discovered page:
  - Fetches the full HTML content (headless Chrome)
  - Extracts text content from HTML (parser/parse\_content.py)
  - Validates content (rejects pages with < 100 chars – likely rendering failures)

▼

### Stage 4: Compare Current vs. Previous Content

- For each page:
  - If NO previous snapshot exists → marks as NEW PAGE
  - If previous snapshot exists → compares old text vs. new text
    - Detects additions (new lines)
    - Detects removals (deleted lines)
    - Detects structural warnings
    - Validates: blocks suspicious changes (>70% shrinkage = rendering failure)
  - Collects all changes into a report

▼

### Stage 5: Save Updated Snapshots & Upload to GCS

- Saves updated/new snapshots to local filesystem inside container
- If GCS is enabled:
  - Uploads ALL local snapshots to GCS bucket
  - Deletes stale GCS files that no longer exist locally (sync)

▼

**Why this stage exists:** The updated snapshots must be persisted to GCS so the **next** run (at 6 PM the same day, or 9 AM the next day) can download them and compare again.

## Stage 6: Send Email Notification (Secret Manager provides password)

- Builds email notification (HTML + plain text) with:
    - | — Summary of changes detected (or "no changes")
    - | — Details of additions/removals per page
    - | — Links to changed pages
  - Reads EMAIL\_PASSWORD from environment variable
    - | (injected by Secret Manager via Cloud Run's --update-secrets)
  - Connects to SMTP server (smtp.office365.com:587)
  - Sends email to configured recipients
- ▼

## Stage 7: Cloud Run Returns Response

- Returns HTTP 200 (success) or HTTP 500 (error) to Cloud Scheduler
  - Container may be kept warm briefly or shut down
- ▼ DONE

## 4. Visual Diagrams

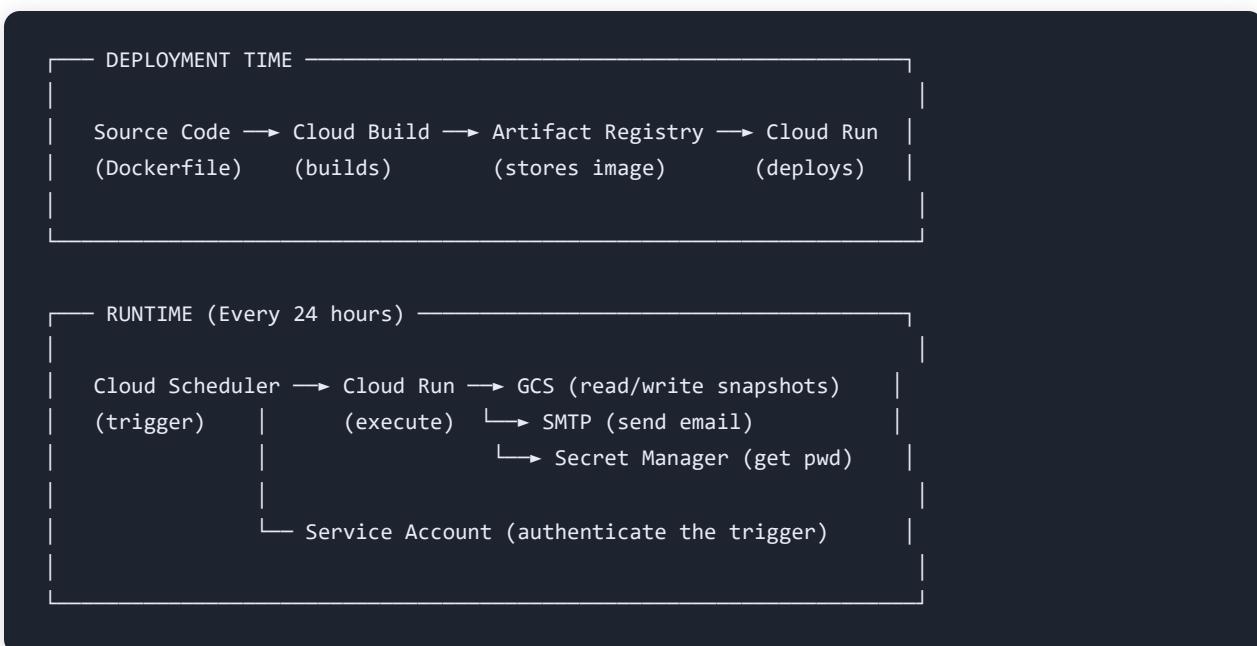
### Deployment Flow (One-Time Setup)



### Runtime Flow (Every Scheduled Run)



## Complete GCP Services Interaction Map



## Summary: Why Each Service Exists

Service	One-Line Purpose
Cloud Build	Builds the Docker image from source code on Google's servers (replaces local <code>docker build</code> + <code>docker push</code> )
Artifact Registry	Stores the built Docker image so Cloud Run can pull it
Cloud Run	Runs the containerized Flask app that performs the monitoring logic
GCS Bucket	Persistent storage for snapshots — because Cloud Run containers are destroyed after each run and lose all local files
Secret Manager	Securely stores the email password — injected into Cloud Run as an env var at runtime
Service Account	Gives Cloud Scheduler an authenticated identity to call the private (no public access) Cloud Run endpoint
Cloud Scheduler	The automated trigger — sends HTTP POST to Cloud Run on a cron schedule (9 AM & 6 PM daily) so the monitoring runs automatically twice a day

**Bottom line:** *Cloud Build + Artifact Registry + Cloud Run = Deployment chain. GCS + Secret Manager + Service Account + Cloud Scheduler = Runtime infrastructure that makes the app work automatically and securely every day.*